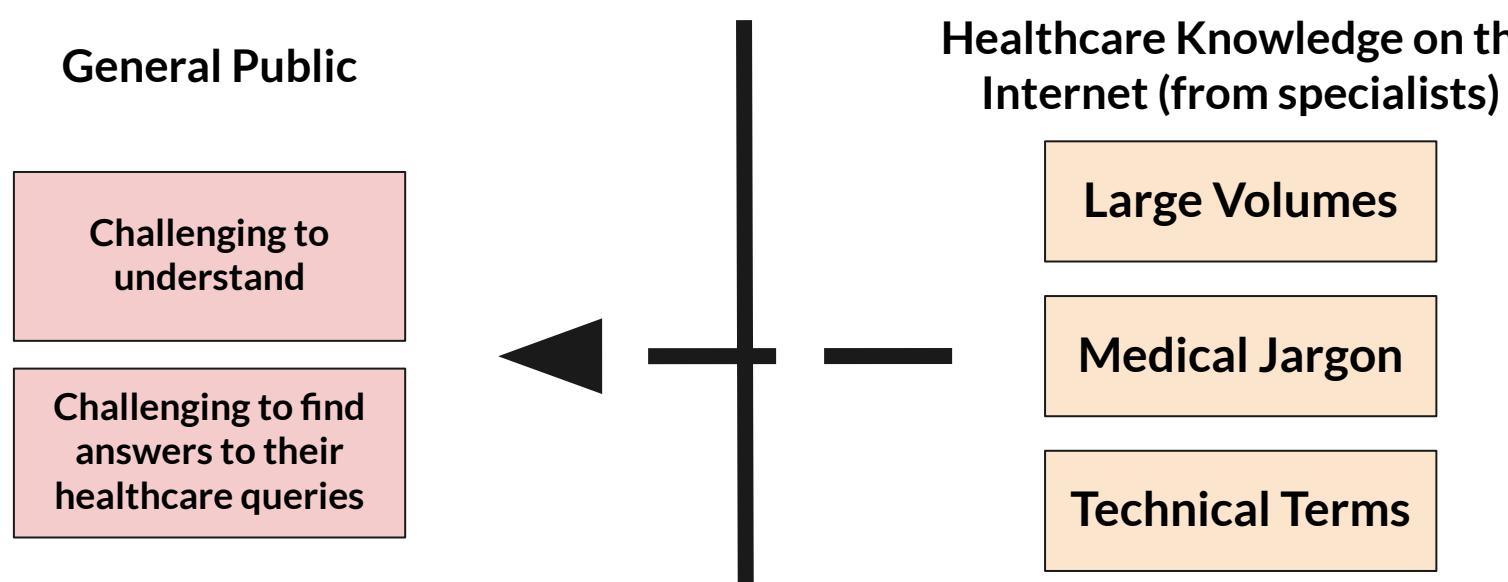



MediLiteQA: Tuning Smaller LMs for Healthcare QA

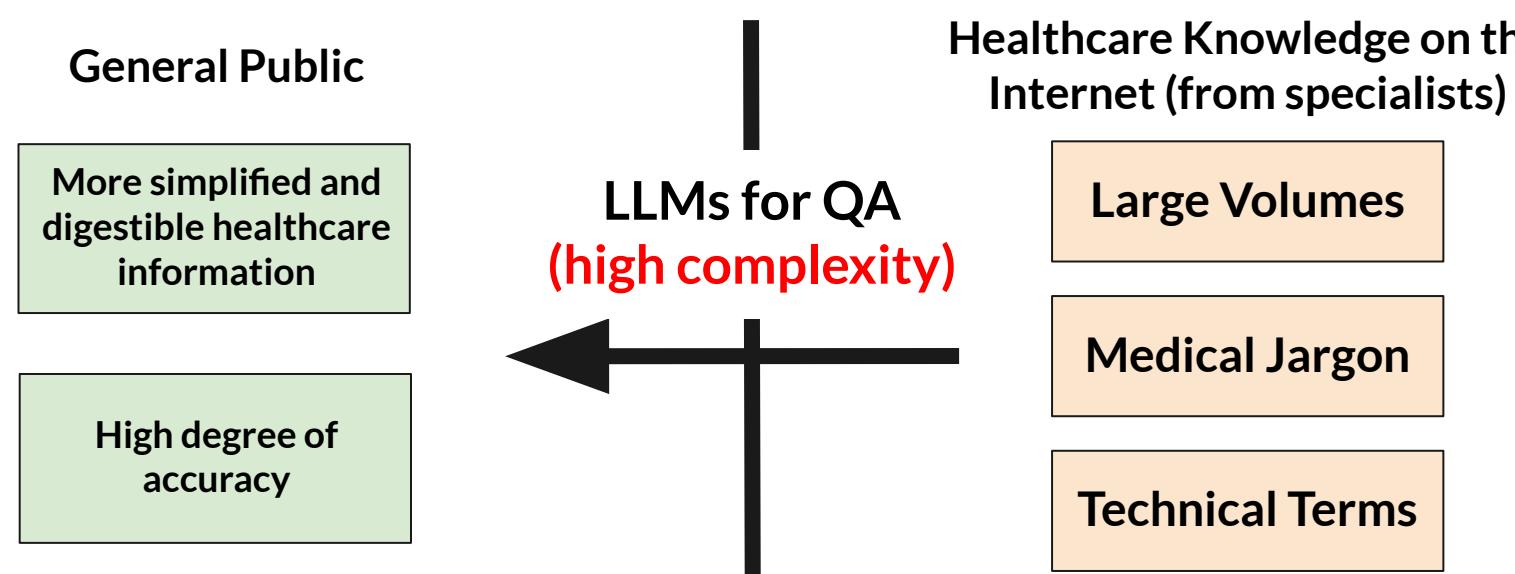
Group 11
Tay Kaiying, Roydon (A0271742E)
Ngu Jia Hao (A0272789H)
Chia Bing Xuan (A0259419R)

1. Introduction

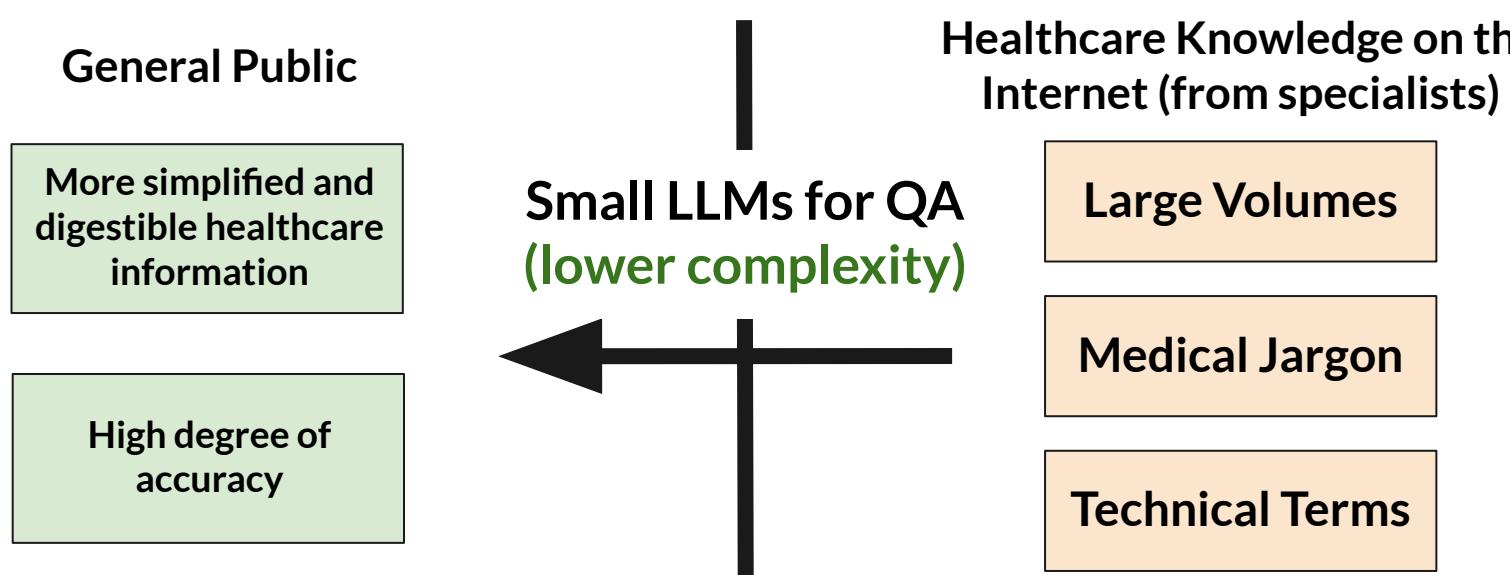
In healthcare, there is a knowledge gap between the layperson and the domain experts



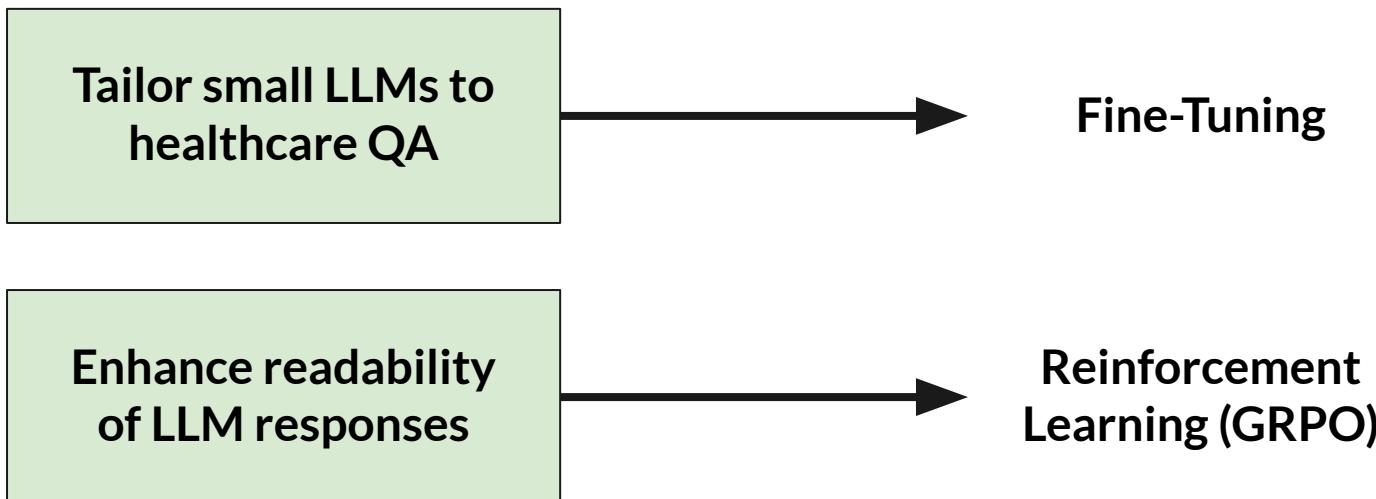
LLMs have bridged this mismatch, but they are expensive to train and store



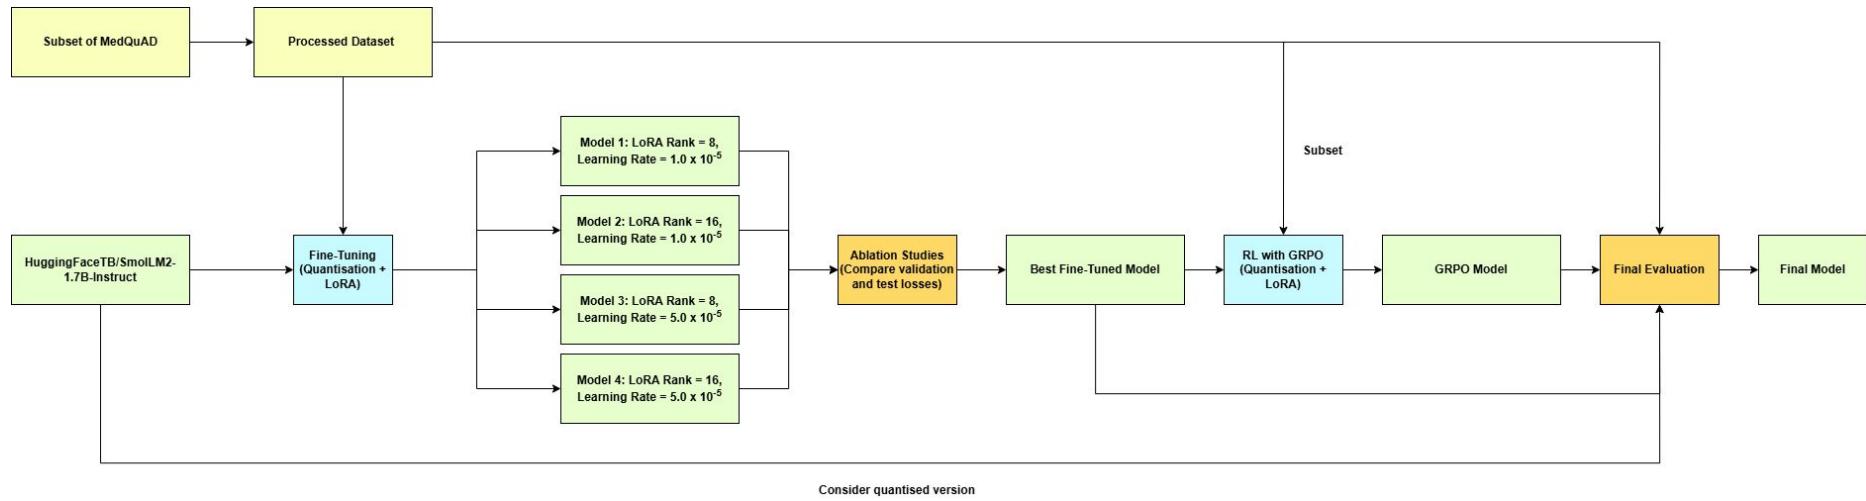
We aim to tune smaller LLMs for healthcare QA



Project Scope: Improving the performance of small LLMs in healthcare QA



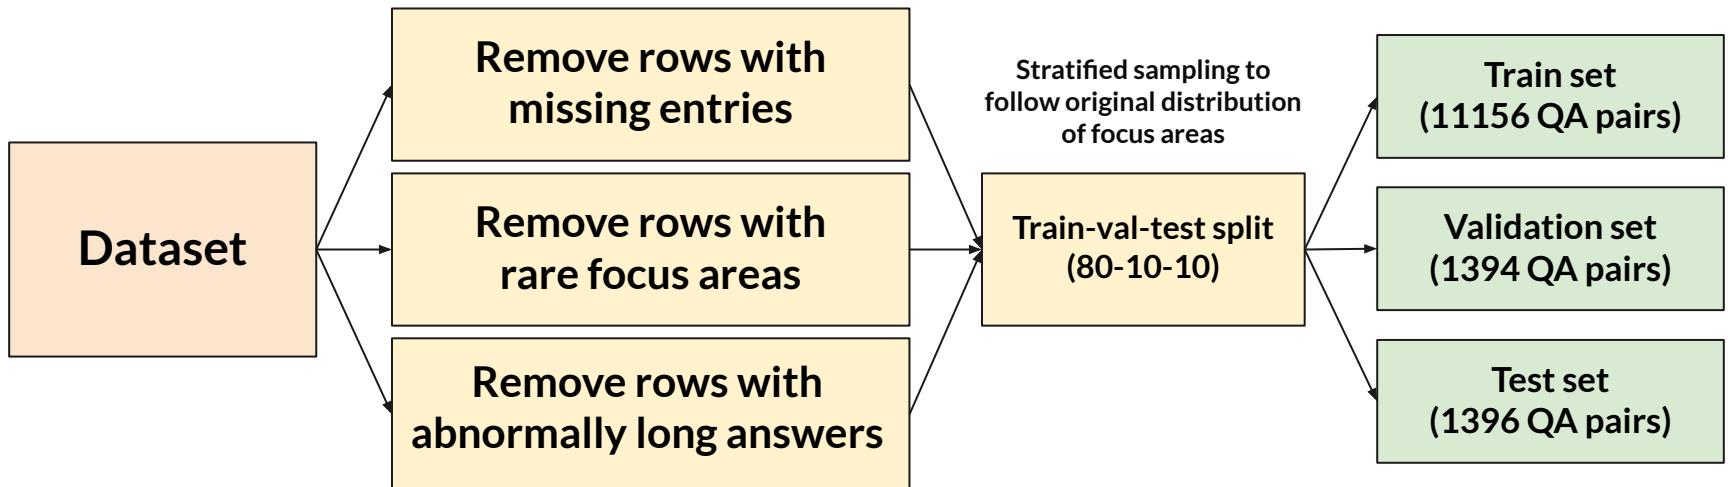
Workflow



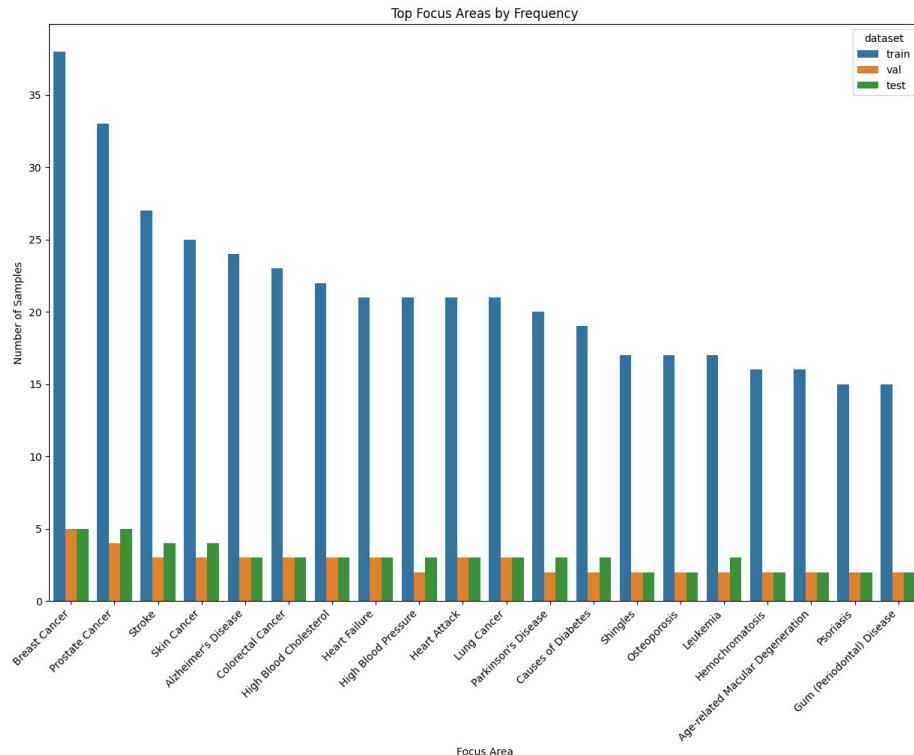
2. Data Cleaning and Processing

Dataset and Data Processing

Subset of Medical Question Answering Dataset (MedQuAD): 12 source websites, 14984 unique question-answer pairs, 5127 focus areas

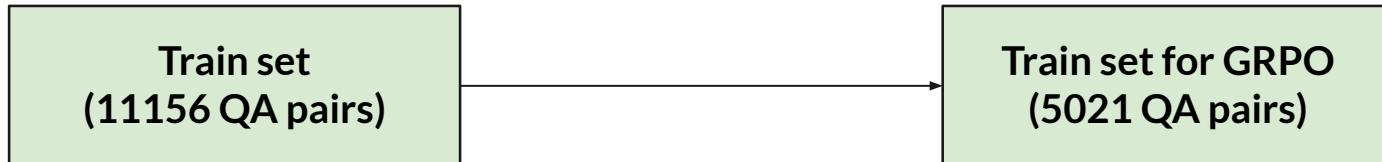


Focus Area Composition of Train-Val-Test Split



Dataset and Data Processing (GRPO)

Owing to computational constraints, we are unable to apply reinforcement learning on the entire training set. We shall use a subset of the training set for GRPO



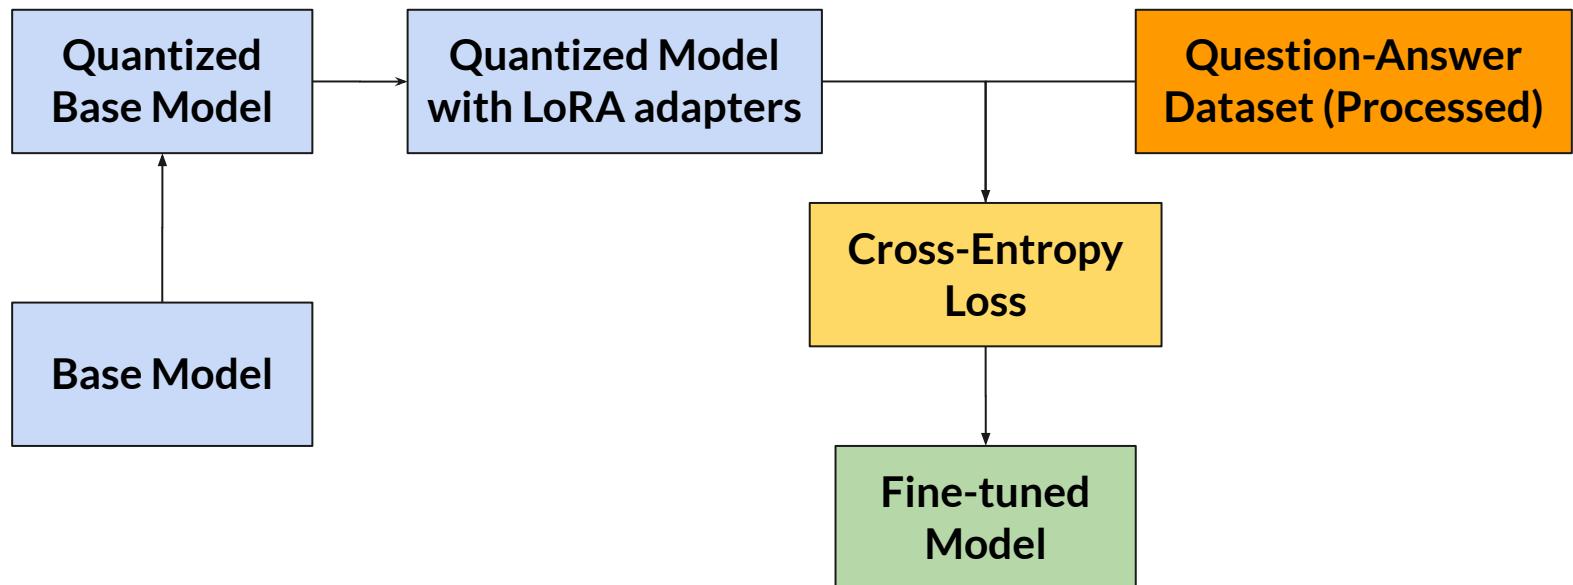
3. Supervised Fine-Tuning

Base Model chosen for fine-tuning

- Base Language model for fine-tuning: HuggingFaceTB/SmoILM2-1.7B-Instruct
- We train this small(er) language model for closed-book QA in the healthcare domain as they are computationally less expensive and can be run on a wider range of devices.
- We aim to improve the domain knowledge of this model in the healthcare domain through fine-tuning.



Fine-tuning model for closed-book QA



Ablation Studies

Model	LoRA Rank	Learning Rate	Validation Loss	Test Loss
1	8	1e-5	1.40	1.40
2	16	1e-5	1.42	1.42
3	8	5e-5	1.13	1.14
4	16	5e-5	1.14	1.14

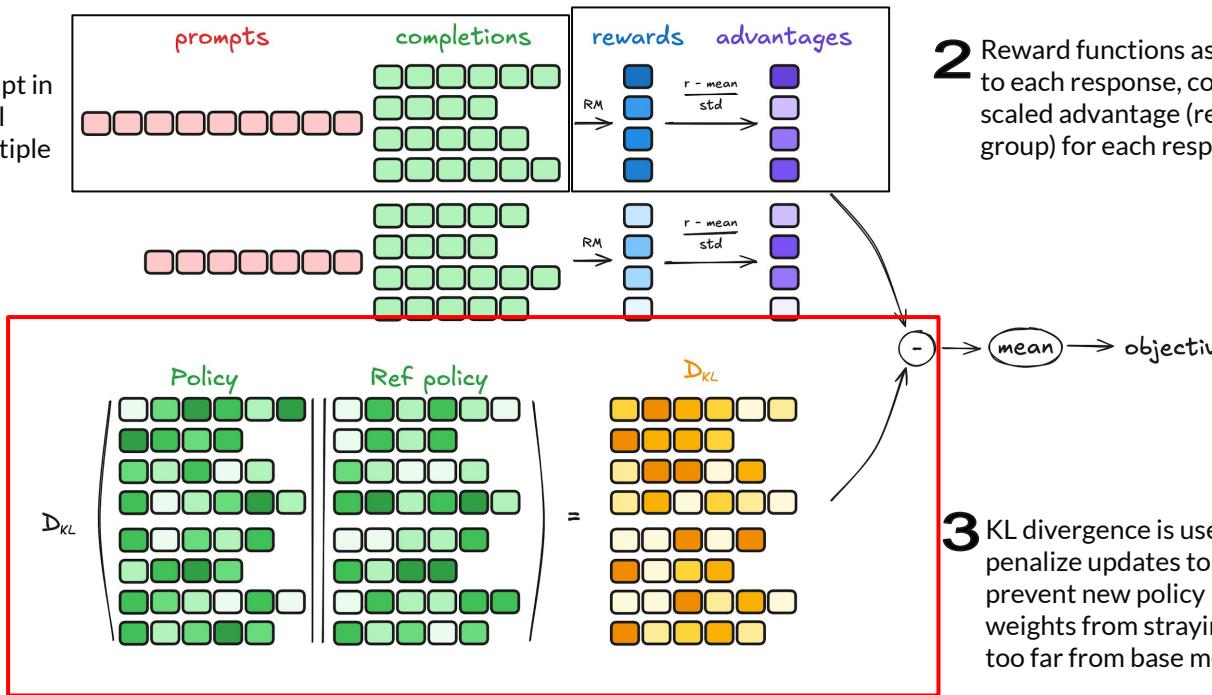
4. Reinforcement Learning (GRPO)

Reinforcement Learning with GRPO

- Why Reinforcement Learning?
 - Want the model to output responses that are easy to read and understand
- Why GRPO?
 - **Better than RL alternatives:** Eliminates the need for a value model (PPO)
Also does not require a curated human preference dataset (DPO)

Reinforcement Learning with GRPO

- 1 For each prompt in dataset, model generates multiple responses (a group)



- 2 Reward functions assign scores to each response, compute the scaled advantage (relative to group) for each response

- 3 KL divergence is used to penalize updates to prevent new policy weights from straying too far from base model

GRPO Objective Function

$$\mathcal{J}_{GRPO}(\theta) = \mathbb{E}[q \sim P(Q), \{o_i\}_{i=1}^G \sim \pi_{\theta_{old}}(O|q)]$$

$$\frac{1}{G} \sum_{i=1}^G \frac{1}{|o_i|} \sum_{t=1}^{|o_i|} \left\{ \min \left[\frac{\pi_\theta(o_{i,t}|q, o_{i,<t})}{\pi_{\theta_{old}}(o_{i,t}|q, o_{i,<t})} \hat{A}_{i,t}, \text{clip} \left(\frac{\pi_\theta(o_{i,t}|q, o_{i,<t})}{\pi_{\theta_{old}}(o_{i,t}|q, o_{i,<t})}, 1 - \varepsilon, 1 + \varepsilon \right) \hat{A}_{i,t} \right] - \beta \mathbb{D}_{KL} [\pi_\theta || \pi_{ref}] \right\}$$



Reward functions chosen

Flesch Reading Ease Score: Readability scores using Average Sentence Length (ASL) and Average Syllables per Word (ASW)

$$\text{Flesch Reading Ease} = 206.835 - (1.015 \times \text{ASL}) - (84.6 \times \text{ASW})$$

Soft Length-Based Punishment

Maximum reward for ideal length ($30 \leq \text{word length} \leq 300$)
Soft decay of reward towards -1 if the length of responses
deviates on either side of ideal length bounds



Reward functions chosen

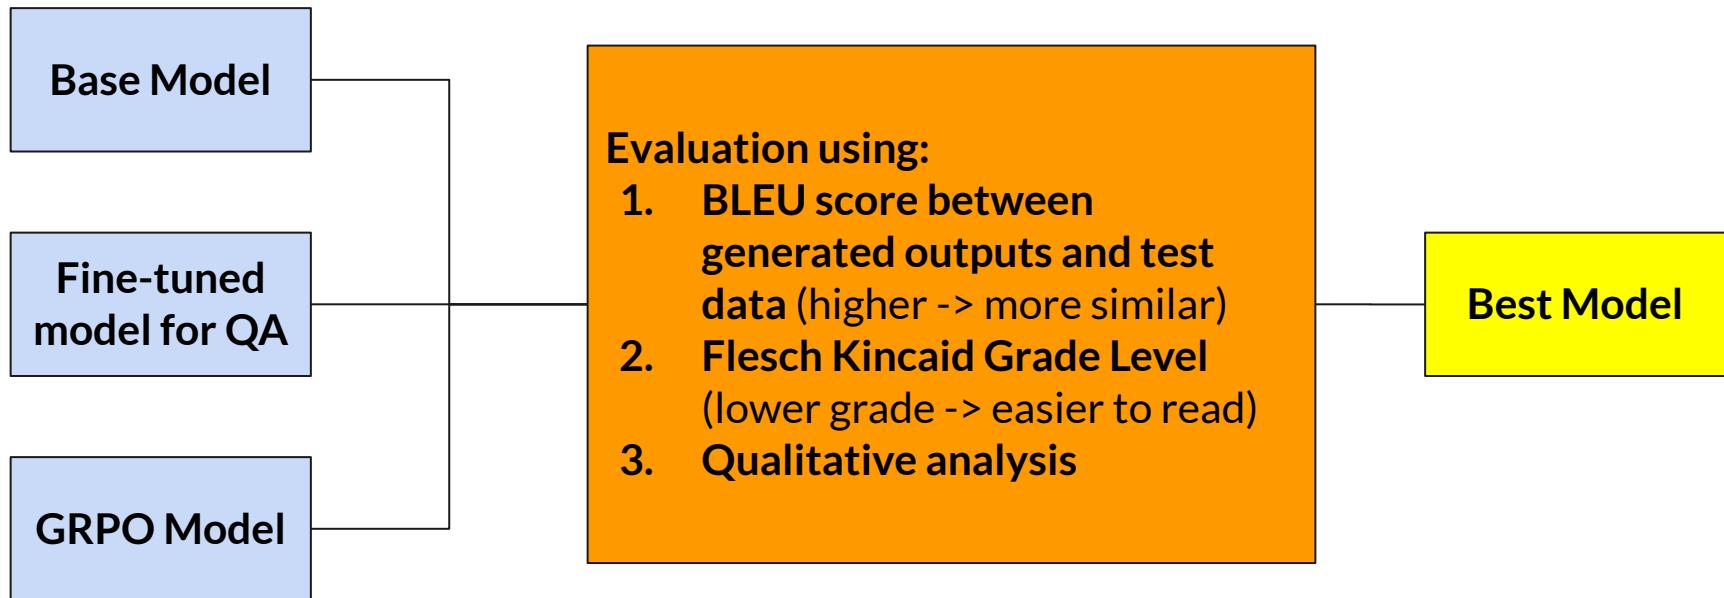
Flesch Reading Ease Score: To get model to discover ways to make its responses more easy to read and understand (e.g. shorter sentences, simpler synonyms)

Soft Length-Based Punishment: To guide model to keep responses within an ideal length

5. Evaluation Results



Evaluation





Evaluation Results

	BLEU	Flesch Kincaid Readability Grade
Base Instruct Model	0.04900	12.3
Fine-tuned Model	0.02044	11.6
GRPO Trained Model	0.02777	10.7

Fine-tuned models still make mistakes

Question: What genetic mutation causes the keratitis-ichthyosis-deafness (KID) syndrome?

["KID syndrome is caused by mutations in the GJB2 gene. This gene provides instructions for making a protein called gap junction beta 2, more commonly known as connexin 26. Connexin 26 is a member of the connexin protein family."]

Answer

The KID syndrome is caused by mutations in the SLC26A4 gene. The SLC26A4 gene provides instructions for making a protein that plays a role in the formation of the skin and the formation of the lens of the eye. The lens is located behind the iris (the colored part of the eye).

Fine-tuned

The KID syndrome is caused by mutations in the SLC26A4 gene. The SLC26A4 gene provides instructions for making a protein that plays a role in the formation of the skin and the formation of the lens of the eye. The lens is located behind the iris (the colored part of the eye).

GRPO

Fine-tuned models got confused when conditions have similar symptoms

(gave another gene that also causes deafness but not the one that causes KID syndrome)



Question: How many people are affected by Bietti crystalline dystrophy (BCD)?

["Bietti crystalline dystrophy (BCD) has been estimated to occur in 1 in 67,000 people. It is more common in people of East Asian descent, especially those of Chinese and Japanese background. Researchers suggest that Bietti crystalline dystrophy may be underdiagnosed because its symptoms are similar to those of other eye disorders that progressively damage the retina."]

Answer

Bietti crystalline dystrophy (BCD) is a disorder characterized by the development of calcium deposits (enlarging crystals) within the lens of the eye. It is an inherited condition that affects males more commonly than females. The prevalence of BCD is unknown.

Fine-tuned

Bietti crystalline dystrophy (BCD) is a rare, inherited eye disorder in which the lens of the eye becomes yellow. The lens is the clear, dome-shaped structure in the front of the eye that helps focus light onto the retina. BCD is also associated with progressive cloudiness of the cornea, the clear, dome-shaped tissue covering the front of the eye. BCD begins during childhood and its symptoms progressively worsen over time. Some people with BCD may also develop cataracts. There is no cure for BCD.

GRPO

Fine-tuned models failed to learn some knowledge from the dataset (did not know the prevalence of BCD although present in data)

Dataset has insufficient training samples for imparting knowledge



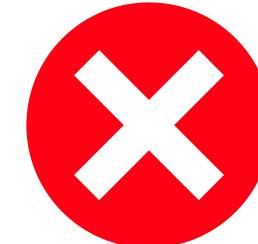
Good and Bad examples from GRPO Training

Fine-tuned	GRPO
<p>Most cases of Diamond-Blackfan anemia are not inherited. They occur when cells in the bone marrow are damaged early in life. This damage is not the result of inherited (genetic) changes. Instead, it is due to a combination of factors, including some unknown genetic (chromosomal) changes.</p>	<p>Most cases of Diamond-Blackfan anemia are not inherited. They occur when cells in the bone marrow are damaged early in life. This damage is not the result of inherited (genetic) changes. Instead, it is due to a combination of factors, including some that are present at birth (inherited).</p>

In successful cases, GRPO trains model to improve word choices while keeping meaning the same
(genetic, chromosomal -> inherited)



However, GRPO caused model to hallucinate facts sometimes.
(hookworm infection should be unrelated to diapers)



GRPO works, but factuality may be affected

Possible Future Work

- **Synthetic Data Creation:** Large LLM can be used to generate quality synthetic data to provide more training samples for smaller LLM (knowledge distillation)
- **Unifying similar training samples:** Dataset contains varied answers for similar questions. These responses could be concatenated and summarized by teacher LLM -> reduces variety of answers for similar / same questions and help model learn better
- **Explore other reward models:** Using LLM-as-Judge for rewarding correctness, reward professional term exclusion

6. Q and A

References

- Abacha, A. B., & Demner-Fushman, D. (2019). A Question-Entailment Approach to Question Answering. *BMC Bioinformatics*, 20(511).
<https://doi.org/10.1186/s12859-019-3119-4>
- Allal, L. B., Lozhkov, A., Bakouch, E., Blázquez, G. M., Penedo, G., Tunstall, L., Marafioti, A., Kydlíček, H., Lajarín, A. P., Srivastav, V., Lochner, J., Fahlgren, C., Nguyen, X.-S., Fourrier, C., Burtenshaw, B., Larcher, H., Zhao, H., Zakka, C., Morlon, M., ... Wolf, T. (2025). SmolLM2: When Smol Goes Big – Data-Centric Training of a Small Language Model. <https://arxiv.org/html/2502.02737v1>
- Cao, Y., Zhao, H., Cheng, Y., Shu, T., Chen, Y., Liu, G., Liang, G., Zhao, J., Yan, J., & Li, Y. (2024). Survey on Large Language Model-Enhanced Reinforcement Learning: Concept, Taxonomy, and Methods. *IEEE Transactions on Neural Networks and Learning Systems*, 36(6), 9737 – 9757.
<https://arxiv.org/abs/2404.00282>
- Dettmers, T., Pagnoni, A., Holtzman, A., & Zettlemoyer, L. (2023). QLoRA: Efficient Finetuning of Quantized LLMs. *Proceedings of the 37th International Conference on Neural Information Processing Systems*, 10088 – 10115. <https://doi.org/10.48550/arXiv.2305.14314>
- Guo, D., Yang, D., Zhang, H., Song, J., Zhang, R., Xu, R., Zhu, Q., Ma, S., Wang, P., Bi, X., Zhang, X., Yu, X., Wu, Y., Wu, Z. F., Gou, Z., Shao, Z., Li, Z., Gao, Z., Liu, A., ... Zhang, Z. (2025). Deepseek-R1: Incentivizing reasoning capability in LLMS via reinforcement learning. [arXiv.org](https://arxiv.org/abs/2501.12948).
<https://arxiv.org/abs/2501.12948>
- Hu, E. J., Shen, Y., Wallis, P., Allen-Zhu, Z., Li, Y., Wang, S., Wang, L., & Chen, W. (2021). LoRA: Low-Rank Adaptation of Large Language Models. *The Tenth International Conference on Learning Representations*. <https://doi.org/10.48550/arXiv.2106.09685>
- Kincaid, J. P., Fishburne, R. P., Jr., Rogers, R. L., & Chissom, B. S. (1975). *Derivation of new readability formulas (Automated Readability Index, Fog Count and Flesch Reading Ease Formula) for Navy enlisted personnel*. Institute for Simulation and Training.
<https://stars.library.ucf.edu/cgi/viewcontent.cgi?article=1055&context=istlibrary>

References

- Leo, C. (2024, July 16). *The Math Behind Multi-Head Attention in Transformers*. Towards Data Science. <https://towardsdatascience.com/the-math-behind-multi-head-attention-in-transformers-c26cba15f625/>
- Miller, J. (2024). Healthcare NLP: LLMs, Transformers, Datasets [Data set]. Kaggle. <https://www.kaggle.com/datasets/jpmiller/layoutlm>
- Qiu, W., Huang, T., Rullo, R., Kuang, Y., Maatouk, A., Ramos, S. R., & Ying, R. (2025). REPHQA: Evaluating Readability of large Language Models in Public Health Question Answering. *arXiv* (Cornell University). <https://doi.org/10.48550/arxiv.2509.16360>
- Schoeninger, G. (2025, February 11). *Why GRPO is Important and How it Works*. Oxen.ai. <https://ghost.oxen.ai/why-grpo-is-important-and-how-it-works/>
- Rafailov, R., Sharma, A., Mitchell, E., Ermon, S., Manning, C. D., & Finn, C. (2023). Direct Preference Optimization: Your Language Model is Secretly a Reward Model. *arXiv.org*. <https://arxiv.org/abs/2305.18290>
- Schulman, J., Wolski, F., Dhariwal, P., Radford, A., & Klimov, O. (2017). Proximal Policy Optimization Algorithms. *arXiv preprint arXiv:1707.06347*. <https://doi.org/10.48550/arXiv.1707.06347>
- Shao, Z., Wang, P., Zhu, Q., Xu, R., Song, J., Bi, X., Zhang, H., Zhang, M., Li, Y. K., Wu, Y., & Guo, D. (2024). DeepSeekMath: Pushing the limits of mathematical reasoning in open language models. *arXiv preprint arXiv:2402.03300v3*. <https://doi.org/10.48550/arXiv.2402.03300>
- Vaswani, A., Shazeer, N., Parmar, N., Uszkoreit, J., Jones, L., Gomez, A. N., Kaiser, Ł., & Polosukhin, I. (2017). Attention Is All You Need. *Advances in Neural Information Processing Systems*, 30, 5998 – 6008. <https://arxiv.org/abs/1706.03762>
- Yu, Q., Zhang, Z., Zhu, R., Yuan, Y., Zuo, X., Yue, Y., Dai, W., Fan, T., Liu, G., Liu, L., Liu, X., Lin, H., Lin, Z., Ma, B., Sheng, G., Tong, Y., Zhang, C., Zhang, W., Zhu, H., ... Wang, M. (2025). DAPO: An open-source LLM reinforcement learning system at scale. *arXiv preprint arXiv:2503.14476*. <https://doi.org/10.48550/arXiv.2503.14476>

References

- Zhang, Y. (2025, February 7). DeepSeek-R1 Dissection: Understanding PPO & GRPO Without Any Prior Reinforcement Learning Knowledge. *Hugging Face.* <https://huggingface.co/blog/NormalUhr/grpo>